

### Question-1:

#### Download the Dataset

#### Solution:

```
from google.colab
import drive
drive.mount('/content/drive')
```

```
# .....#
# .....#
```

#### Download the Dataset

```
In [2]: from google.colab import drive
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force\_remount=True).

## Question-2:

## Image Augmentation

### Solution :

#### Image Augmentation

```
In [3]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from matplotlib import style
import seaborn as sns
import cv2
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import os
import PIL
import random
import cv2
from tensorflow.keras import layers, models
import tensorflow as tf
import pandas as pd
from sklearn.model_selection import train_test_split
import seaborn as sns
import pickle
import zipfile
tf.__version__
```

Out[3]: '2.8.2'

```
In [4]: !ls
```

drive sample\_data

```
In [5]: try:
    tpu = tf.distribute.cluster_resolver.TPUClusterResolver()
    print('Device:', tpu.master())
    tf.config.experimental_connect_to_cluster(tpu)
    tf.tpu.experimental.initialize_tpu_system(tpu)
    strategy = tf.distribute.experimental.TPUStrategy(tpu)
except:
    strategy = tf.distribute.get_strategy()
print('Number of replicas:', strategy.num_replicas_in_sync)
```

Number of replicas: 1

```
In [6]: AUTOTUNE = tf.data.experimental.AUTOTUNE
batch_size = 32
IMAGE_SIZE = [128, 128]
EPOCHS = 25
```

```
In [7]: image = cv2.imread(r'/content/drive/MyDrive/Flowers-Dataset/flowers/daisy/100080576_f52e8ee070_n.jpg')
```


```
In [8]: print(image.shape)
```

(263, 320, 3)

```
In [9]: imgplot = plt.imshow(image)
plt.show()
```

0





```
GCS_PATH= "/content/drive/MyDrive/Flowers-Dataset/flowers™
```

```
CLASS_NAMES = up.array([str(th.strings.split(item, os.path. sep) [-1]).numpy() [2: -1]
for item in tf.io.gfile.glob(str(GCS_PATH + "*/"))])
CLASS_NAMES
```

```
array(['daisy', 'rose', 'dandelion', 'sunflower', 'tulip'], dtype='<U9')
```

```
files_count = []
for i, f in enumerate(CLASS_NAMES):
    folder_path = os.path.join(GCS_PATH, f)
    for path in os.listdir(os.path.join(folder_path, f)):
        files_count.append([f'({f})'.format(folder_path, path), f, i])
flowers_df = pd.DataFrame(files_count, columns=['filepath', 'class_name', 'label'])
flowers_df.head()
```

	filepath	class_name	label
0	/content/drive/MyDrive/Flowers-Dataset/flowers...	daisy	0
1	/content/drive/MyDrive/flowers-Dataset/flowers...	daisy	0
2	/content/drive/MyDrive/Flowers-Dataset/flowers...	daisy	0
3	/content/drive/MyDrive/Flowers-Dataset/flowers...	daisy	0
4	/content/drive/MyDrive/Flowers-Dataset/flowers...	daisy	0

```
flowers_df.class_name.value_counts()
```

```
dandelion    1852
tulip         984
rose          784
daisy         764
sunflower    733
```

```
have: class_name, dtype: object
```

```
quantidade_por_class = 500
```

```
flowers_df = pd.concat([flowers_df[flowers_df['class_name'] == i] [: quantidade_por_class] for i in CLASS_NAMES])
```


```
flowers_df.class_name.value_counts()
```

```
daisy        500
```

```
dandelion    500
sunflower    500
tulip         500
```

```
sunflower    500  
tulip        500  
Name: class_name, dtype: int64
```

```
In [15]: image = cv2.imread(flowers_df.filepath[100])  
imgplot = plt.imshow(image)  
plt.show()
```



#### Create Model

```
In [16]: X = flowers_df['filepath']  
y = flowers_df['label']  
  
x_train, x_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=101)
```

```
In [17]:
```

### Question-3:

### Create Model

#### Solution :

```
In [16]: X = flowers_df['filepath']
y = flowers_df['label']

x_train, x_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=101)
```

```
In [17]: x_train_tensor = tf.convert_to_tensor(x_train.values, dtype=tf.string)
y_train_tensor = tf.convert_to_tensor(y_train.values)

x_test_tensor = tf.convert_to_tensor(x_test.values, dtype=tf.string)
y_test_tensor = tf.convert_to_tensor(y_test.values)
```



```
In [18]: train_data = tf.data.Dataset.from_tensor_slices((x_train_tensor, y_train_tensor))
test_data = tf.data.Dataset.from_tensor_slices((x_test_tensor, y_test_tensor))
```

```
In [19]: def map_fn(path, label):
        image = tf.image.decode_jpeg(tf.io.read_file(path))

        return image, label

#apply the function
train_data_img = train_data.map(map_fn)
test_data_img = test_data.map(map_fn)
```

```
In [20]: fig, ax = plt.subplots(1,2, figsize = (15,5))
for i,l in train_data_img.take(1):
    ax[0].set_title('Image dataset to train');
    ax[0].imshow(i);
for i,l in test_data_img.take(1):
    ax[1].set_title('Image dataset to test');
    ax[1].imshow(i);
```



```

def preprocess(image):
    """
    Returns a image that is reshaped and normalized

    image = I.E. cast image, I.F. Float32)

    image = tf.image.resize(image, IMAGE_SIZE)

    return image, Zabel



train_data_norm = train_data_img.map(preprocessing)

```

```

fig, ax = plt.subplots(1, 2, figsize = (15, 5))
for i, l in train_data_norm.take(1):
    ax[0].set_title('Image dataset to train');
    ax[0].imshow(i);
for i, l in test_data_norm.take(1):
    ax[1].set_title('Image dataset to test');
    ax[1].imshow(i);

```



```

train_batches = train_data_norm.batch(batch_size)
test_batches = test_data_norm.batch(batch_size)

for i, l in train_batches.take(1):
    print('Train Data Shape', i.shape)
for i, l in test_batches.take(1):
    print('Test Data Shape', i.shape)

```

Train Data Shape (32, 128, 128, 3)  
 Test Data Shape (32, 128, 128, 3)

#### Question-4 :

#### Add Layers (Convolution,MaxPooling,Flatten,Dense-(Hidden Layers),Output)

##### Solution :

##### Add Layers (Convolution,MaxPooling,Flatten,Dense-(Hidden Layers),Output)

```
In [24]: LeNet = models.Sequential()
LeNet.add(layers.Conv2D(6, (5,5), activation = 'relu', input_shape = (128, 128, 3)))
LeNet.add(layers.MaxPooling2D())
LeNet.add(layers.Conv2D(16, (5,5), activation = 'relu'))
LeNet.add(layers.MaxPooling2D())
LeNet.add(layers.Flatten())
LeNet.add(layers.Dense(255, activation='relu'))
LeNet.add(layers.Dropout(0.2))
LeNet.add(layers.Dense(124, activation='relu'))
LeNet.add(layers.Dropout(0.2))
LeNet.add(layers.Dense(84, activation='relu'))
LeNet.add(layers.Dense(43, activation='sigmoid'))
LeNet.summary()
```

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 124, 124, 6)	456
max_pooling2d (MaxPooling2D)	(None, 62, 62, 6)	0
conv2d_1 (Conv2D)	(None, 58, 58, 16)	2416
max_pooling2d_1 (MaxPooling2D)	(None, 29, 29, 16)	0
flatten (Flatten)	(None, 13456)	0
dense (Dense)	(None, 255)	3431535
dropout (Dropout)	(None, 255)	0
dense_1 (Dense)	(None, 124)	31744
dropout_1 (Dropout)	(None, 124)	0
dense_2 (Dense)	(None, 84)	10500
dense_3 (Dense)	(None, 43)	3655
=====		
Total params: 3,480,306		
Trainable params: 3,480,306		
Non-trainable params: 0		

#### Question-5:

#### Compile The Model

##### Solution :

##### Compile The Model

```
In [25]: LeNet.compile(optimizer='Adam', loss='sparse_categorical_crossentropy',
metrics = ['accuracy'])
```

## Question-6: Fit The Model


Solution :

### Fit The Model

```
In [26]: history = LeNet.fit(train_batches, epochs=10, batch_size = 16, validation_data=(test_batches))
```

```
Epoch 1/10
55/55 [=====] - 130s 2s/step - loss: 1.7673 - accuracy: 0.2943 - val_loss: 1.3046 - val_accuracy: 0.4560
Epoch 2/10
55/55 [=====] - 40s 724ms/step - loss: 1.2971 - accuracy: 0.4434 - val_loss: 1.2441 - val_accuracy: 0.4880
Epoch 3/10
55/55 [=====] - 42s 752ms/step - loss: 1.1785 - accuracy: 0.5034 - val_loss: 1.1907 - val_accuracy: 0.5173
Epoch 4/10
55/55 [=====] - 36s 650ms/step - loss: 1.0667 - accuracy: 0.5526 - val_loss: 1.1468 - val_accuracy: 0.5453
Epoch 5/10
55/55 [=====] - 49s 889ms/step - loss: 0.9430 - accuracy: 0.6366 - val_loss: 1.1333 - val_accuracy: 0.5520
Epoch 6/10
55/55 [=====] - 37s 673ms/step - loss: 0.7835 - accuracy: 0.7051 - val_loss: 1.2531 - val_accuracy: 0.5333
Epoch 7/10
55/55 [=====] - 36s 648ms/step - loss: 0.6586 - accuracy: 0.7531 - val_loss: 1.2900 - val_accuracy: 0.5427
Epoch 8/10
55/55 [=====] - 40s 719ms/step - loss: 0.4778 - accuracy: 0.8257 - val_loss: 1.6341 - val_accuracy: 0.5080
Epoch 9/10
55/55 [=====] - 36s 647ms/step - loss: 0.3595 - accuracy: 0.8703 - val_loss: 2.0376 - val_accuracy: 0.4947
Epoch 10/10
55/55 [=====] - 41s 744ms/step - loss: 0.2947 - accuracy: 0.9023 - val_loss: 2.2624 - val_accuracy: 0.4693
```

```
In [31]: plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Model Loss')
plt.ylabel('Loss')
plt.xlabel('Epochs')
plt.legend(['train', 'test'])
plt.show()
```





## Question-7:

### Save the Model

Solution :

#### Save the Model

```
In [32]: from sklearn.neighbors import KNeighborsClassifier as KNN
import numpy as np

# Load dataset
from sklearn.datasets import load_iris
iris = load_iris()

X = iris.data
y = iris.target

# Split dataset into train and test
X_train, X_test, y_train, y_test = \
    train_test_split(X, y, test_size=0.3,
                    random_state=2018)

# import KNeighborsClassifier model
knn = KNN(n_neighbors=3)

# train model
knn.fit(X_train, y_train)
```

```
Out[32]: KNeighborsClassifier(n_neighbors=3)
```

```
In [30]: import pickle
saved_model = pickle.dumps(knn)
knn_from_pickle = pickle.loads(saved_model)
knn_from_pickle.predict(X_test)
```

```
Out[30]: array([0, 1, 1, 1, 0, 1, 2, 1, 2, 0, 0, 2, 2, 2, 0, 2, 2, 0, 1, 1, 1, 0,
        2, 0, 0, 2, 0, 0, 2, 1, 0, 2, 0, 1, 2, 0, 0, 0, 0, 1, 0, 2, 2, 2,
        1])
```

### Question-8:

#### Test The Model

Solution :

#### Test The Model

```
In [27]: import warnings
warnings.filterwarnings('always')
warnings.filterwarnings('ignore')
```

```
In [28]: plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('Model Accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epochs')
plt.legend(['train', 'test'])
plt.show()
```

